

R Notebook

[Code ▼](#)

This is an R Markdown (<http://rmarkdown.rstudio.com>) Notebook. When you execute code within the notebook, the results appear beneath the code.

Try executing this chunk by clicking the *Run* button within the chunk or by placing your cursor inside it and pressing *Cmd+Shift+Enter*.

—SVM classification—

Red dataset

[Hide](#)

```
library(kernlab)
library(e1071)
library(MASS)

set.seed(1)
red <- read.csv('winequality-red.csv', header = TRUE, sep=";")
red <- na.omit(red)
red.quality <- red$quality
red[, -12] <- scale(red[, -12])

# Red dataset
train_red_idx <- sample(nrow(red) * 0.8) # 80-20 train-test split
train.red <- red[train_red_idx,]
train.red$quality <- as.factor(train.red$quality)
train.red.quality <- train.red$quality

test.red <- red[-train_red_idx,]
test.red$quality <- as.factor(test.red$quality)
test.red.quality <- test.red$quality
# test.red <- test.red[, -12]

# # Backward model selection
# library(leaps)
# regfit.bwd=regsubsets(quality~., data=train.red, nvmax=11, method="backward")
# summary(regfit.bwd)
#
# test.mat=model.matrix(quality~., data=test.red) # create an X matrix of test data
# val.errors=rep(NA, 19)
# for(i in 1:19){
#   coefi=coef(regfit.best, id=i)
#   pred=test.mat[, names(coefi)]%*%coefi
#   val.errors[i]=mean((Hitters$Salary[test]-pred)^2)
# }
# val.errors
# which.min(val.errors)
# coef(regfit.best, 10)

set.seed(1)

gamma <- 2^seq(-15, 3, 1)
test_accuracy <- rep(0, length(gamma))
fit_time <- rep(0, length(gamma))
i = 1

start <- proc.time()

for(g in gamma){
  start_it <- proc.time()
  model = svm(quality~., data = train.red, gamma=g)
  fit_time[i] <- proc.time() - start_it
  test_accuracy[i] <- mean(test.red.quality == predict(model, test.red))
  i = i + 1
}
```

[illegible]

Hide

```
print(cv_time_red <- proc.time() - start) # total CV time
```

```
user    system elapsed
3.955   0.095   4.214
```

Hide

```
print(fit_time[which.max(test_accuracy)]) # fit time of best model
```

```
[1] 0.15
```

Hide

```
print(best_gamma_red <- 2^(which.max(test_accuracy) - 16)) # gamma yielding highest test accuracy
```

```
[1] 0.25
```

Hide


```
print(test_acc_red <- max(test_accuracy)) # test accuracy
```

```
[1] 0.628125
```

Hide


```
plot(log2(gamma), test_accuracy, type='b', xlab = 'log_2(gamma)', main='Test Accuracies for different gamma via S  
VM Classification on Red Subset')
```

Test Accuracies for different gamma via SVM Classification on Red Subse



Hide

```
plot(log2(gamma), fit_time, type='b', xlab = 'log_2(gamma)', main='Fit Times for different gamma via SVM Classification on Red Subset')
```



Hide

NA
NA

White dataset

Hide

```
library(kernlab)  
library(e1071)
```

package 'e1071' was built under R version 3.6.2

Hide


```
[1] 1.668
```

Hide

```
print(best_gamma_white <- 2^(which.max(test_accuracy) - 16)) # gamma yielding highest test accuracy
```

```
[1] 0.015625
```

Hide


```
print(test_acc_white <- max(test_accuracy)) # test accuracy
```

```
[1] 0.6020408
```

Hide

```
plot(log2(gamma), test_accuracy, type='b', xlab = 'log_2(gamma)', main='Test Accuracies for different gamma via SVM Classification on White Subset')
```


Test Accuracies for different gamma via SVM Classification on White Subs



Hide

```
plot(log2(gamma), fit_time, type='b', xlab = 'log_2(gamma)', main='Fit Times for different gamma via SVM Classification on White Subset')
```

Fit Times for different gamma via SVM Classification on White Subset



Weighted Test Accuracy

Hide

```
(nrow(test.red) * test_acc_red + nrow(test.white) * test_acc_white) / (nrow(test.red) + nrow(test.white))
```

```
[1] 0.6084615
```

— SVM regression (epsilon-insensitive loss) —

Red dataset

Hide

```
set.seed(1)
red <- read.csv('winequality-red.csv', header = TRUE, sep=";")
red <- na.omit(red)
red.quality <- red$quality
red[, -12] <- scale(red[, -12])

# Red dataset
train_red_idx <- sample(nrow(red) * 0.8) # 80-20 train-test split
train.red <- red[train_red_idx,]
# train.red$quality <- as.factor(train.red$quality)
train.red.quality <- train.red$quality

test.red <- red[-train_red_idx,]
# test.red$quality <- as.factor(test.red$quality)
test.red.quality <- test.red$quality
# test.red <- test.red[, -12]

# # Backward model selection
# library(leaps)
# regfit.bwd=regsubsets(quality~.,data=train.red,nvmax=11,method="backward")
# summary(regfit.bwd)
#
# test.mat=model.matrix(quality~.,data=test.red) # create an X matrix of test data
# val.errors=rep(NA,19)
# for(i in 1:19){
#   coefi=coef(regfit.best,id=i)
#   pred=test.mat[,names(coefi)]%*%coefi
#   val.errors[i]=mean((Hitters$Salary[test]-pred)^2)
# }
# val.errors
# which.min(val.errors)
# coef(regfit.best,10)

set.seed(1)

gamma <- 2^seq(-15, 3, 1)
test_accuracy <- rep(0, length(gamma))
fit_time <- rep(0, length(gamma))
i = 1

start <- proc.time()

for(g in gamma){
  start_it <- proc.time()
  model = svm(quality~., data = train.red, gamma=g, type='eps-regression')
  fit_time[i] <- proc.time() - start_it
  test_accuracy[i] <- mean(abs(test.red.quality - predict(model, test.red)) <= 0.5)
  i = i + 1
}
```

[illegible]

Hide

```
print(cv time red <- proc.time() - start) # total CV time
```

```
user    system elapsed
3.205    0.057    3.267
```

Hide

```
print(fit_time[which.max(test_accuracy)]) # fit time of best model
```

```
[1] 0.128
```

Hide

```
print(best_gamma_red <- 2^(which.max(test_accuracy) - 16)) # gamma yielding highest test accuracy
```

```
[1] 0.001953125
```

Hide


```
print(test_acc_red <- max(test_accuracy)) # test accuracy
```

```
[1] 0.640625
```

Hide

```
plot(log2(gamma), test_accuracy, type='b', xlab = 'log_2(gamma)', main='Test Accuracies for different gamma via SVM Regression on Red Subset')
```


Test Accuracies for different gamma via SVM Regression on Red Subset



Hide

```
plot(log2(gamma), fit_time, type='b', xlab = 'log_2(gamma)', main='Fit Times for different gamma via SVM Regression on Red Subset')
```


Fit Times for different gamma via SVM Regression on Red Subset



White dataset

Hide

```
white <- read.csv('winequality-white.csv', header = TRUE, sep=";")
white <- na.omit(white)
white.quality <- white$quality
white[, -12] <- scale(white[, -12])

set.seed(1)
train_white_idx <- sample(nrow(white) * 0.8) # 80-20 train-test split
train.white <- white[train_white_idx,]
# train.white$quality <- as.factor(train.white$quality)
train.white.quality <- train.white$quality

test.white <- white[-train_white_idx,]
# test.white$quality <- as.factor(test.white$quality)
test.white.quality <- test.white$quality
# test.white <- test.white[, -12]

# # Backward model selection
# library(leaps)
# regfit.bwd=regsubsets(quality~., data=train.white, nvmax=11, method="backward")
# summary(regfit.bwd)
#
# test.mat=model.matrix(quality~., data=test.white) # create an X matrix of test data
# val.errors=rep(NA, 19)
# for(i in 1:19){
#   coefi=coef(regfit.best, id=i)
#   pred=test.mat[, names(coefi)]%*%coefi
#   val.errors[i]=mean((Hitters$Salary[test]-pred)^2)
# }
# val.errors
# which.min(val.errors)
# coef(regfit.best, 10)

set.seed(1)

gamma <- 2^seq(-15, 3, 1)
test_accuracy <- rep(0, length(gamma))
fit_time <- rep(0, length(gamma))
i = 1

start <- proc.time()

for(g in gamma){
  start_it <- proc.time()
  model = svm(quality~., data = train.white, gamma=g, type='eps-regression')
  fit_time[i] <- proc.time() - start_it
  test_accuracy[i] <- mean(abs(test.white.quality - predict(model, test.white)) <= 0.5)
  i = i + 1
}
```

```
number of items to replace is not a multiple of replacement length  
number of items to replace is not a multiple of replacement length  
number of items to replace is not a multiple of replacement length  
number of items to replace is not a multiple of replacement length  
number of i  
tems to replace is not a multiple of replacement length  
number of items to replace is not a multiple of replacemen  
t length  
number of items to replace is not a multiple of replacement length  
number of items to replace is not a multiple of replacement length  
multiple of replacement length  
number of items to replace is not a multiple of replacement length  
number of items to replace is not a multiple of replacement length  
umber of items to replace is not a multiple of replacement length  
number of items to replace is not a multiple of replacement length  
replacement length  
number of items to replace is not a multiple of replacement length  
number of items to replace is not a multiple of replacement length  
is not a multiple of replacement length  
number of items to replace is not a multiple of replacement length  
number of i  
tems to replace is not a multiple of replacement length  
number of items to replace is not a multiple of replacemen  
t length
```

Hide

```
print(cv time white <- proc.time() - start) # total CV time
```

```

user    system elapsed
27.852   0.369   28.544

```

Hide

```
print(fit_time[which.max(test_accuracy)]) # fit time of best model
```

[1] 1.276

Hide

```
print(best_gamma_white <- 2^(which.max(test_accuracy) - 16)) # gamma yielding highest test accuracy
```

```
[1] 0.125
```

Hide


```
print(test_acc_white <- max(test_accuracy)) # test accuracy
```

```
[1] 0.5836735
```

Hide

```
plot(log2(gamma), test_accuracy, type='b', xlab = 'log_2(gamma)', main='Test Accuracies for different gamma via SVM Regression on White Subset')
```


Test Accuracies for different gamma via SVM Regression on White Subse



Hide

```
plot(log2(gamma), fit_time, type='b', xlab = 'log_2(gamma)', main='Fit Times for different gamma via SVM Regression on White Subset')
```

Fit Times for different gamma via SVM Regression on White Subset



Weighted Test Accuracy

Hide

```
(nrow(test.red) * test_acc_red + nrow(test.white) * test_acc_white) / (nrow(test.red) + nrow(test.white))
```

```
[1] 0.5976923
```

When you save the notebook, an HTML file containing the code and output will be saved alongside it (click the *Preview* button or press *Cmd+Shift+K* to preview the HTML file).

The preview shows you a rendered HTML copy of the contents of the editor. Consequently, unlike *Knit*, *Preview* does not run any R code chunks. Instead, the output of the chunk when it was last run in the editor is displayed.

When you save the notebook, an HTML file containing the code and output will be saved alongside it (click the *Preview* button or press *Cmd+Shift+K* to preview the HTML file).

The preview shows you a rendered HTML copy of the contents of the editor. Consequently, unlike *Knit*, *Preview* does not run any R code chunks. Instead, the output of the chunk when it was last run in the editor is displayed.